Briefing | A test of faith

Anthropic's astonishing commercial success makes it a target

The cultlike firm is battling with the government for control of superpowered AI

Save Share

Jun 18th 2026 | SAN FRANCISCO | 13 min read

Listen to this story

0:00 / 0:00

VIDEO: TIMO LENZEN

THE PARADOX of religious authority is that, to believers, there is nothing more important, yet spiritual leaders typically wield little practical power. Indeed, it is often the refusal to sully themselves with the quandaries and compromises of everyday life that confers moral authority on holy men. That, in turn, makes it both easy for worldly leaders to subdue them in the short run (they have no weapons) and difficult in the longer term (their persecution only increases their moral standing).

This dynamic, although familiar to medieval prelates and princes, may not be as well understood by Dario Amodei, the supreme pontiff of Anthropic, one of America's leading artificial-intelligence firms, and Donald Trump, America's imperious president. In recent days, for the second time in four months, underlings of Mr Trump have interfered peremptorily in the running of Anthropic. They ordered it to prevent any non-Americans, including those on its own staff, from accessing its latest models, Fable 5 and Mythos 5, on grounds of national security.

Whereas the Trump administration sees AI as a potent tool to strengthen America's position in the world, Anthropic sees it as a force too powerful to be conferred without safeguards on impetuous laypeople. The result is a titanic clash between church and state—one which not only threatens to interfere with Anthropic's blockbuster IPO, but which will also determine who has ultimate control over the world's mightiest technology.

Anthropic's headquarters in San Francisco have a monastic feel. Employees are on a mission to create AI so powerful that insiders liken their task to the mystery of creation. Much of their talk is about ethics and values. When Pope Leo XIV issued an encyclical on AI last month ("*Magnifica humanitas*"), a cofounder of Anthropic was at his side. Mr Amodei, meanwhile, writes frequent encyclicals of his own on how to build AI safely and responsibly. Like the pope's, they dwell on the risks—but they also address the rewards. His latest ("Policy on the AI Exponential") was published on June 10th. It urged the American government to move quickly to establish "serious and binding regulation" of AI.

Superpowered simony

This piety has not hindered Anthropic's business—far from it. The firm was founded only in 2021, shortly after Mr Amodei, his sister Daniela and six others left OpenAI, then the leading AI lab, convinced that it was not serious enough about safety. For more than two years, Anthropic was in OpenAI's shadow. Although it released its own chatbot, Claude, shortly after OpenAI released ChatGPT, it was OpenAI's offering that won most attention.

Anthropic's slow start forced it to focus, however. Rather than doubling down on chatbots, which have many non-paying users and require subscriptions or advertising to make profits, it took a narrower bet on corporate customers. Whereas OpenAI dabbled in lots of consumer products, Anthropic concentrated on creating business tools, such as coding assistants, that generated more durable income and have become good enough to be able to do much of the work of improving themselves.

It has been a winning wager. Software developers have flocked to Claude Code, Anthropic's year-old coding agent, because it is capable of tasks that would take human coders hours (or days) and its work can easily be checked. Business customers like Anthropic's emphasis on safeguards, which tallies with their concerns about rogue AI. Its focus on coding enabled its models to get an early lead on competitors such as OpenAI and Google on benchmarks measuring software-engineering capabilities. Its coding prowess, requiring step-by-step reasoning and tool use, gives it an advantage in creating AI agents to work alongside humans.

OpenAI has pivoted towards coding this year, and its latest models are considered as good as Anthropic's on some tasks. It, too, stands to benefit from the same agentic tailwinds. But four-fifths of Anthropic's revenue comes from business customers. At OpenAI, business clients generate only about 40% of income. That difference gives Anthropic a greater ability to make money.

This month both Anthropic and OpenAI filed for IPOs. Each is likely to be valued at around $1trn. Inevitably, their relative performance is being closely scrutinised. Neither has yet made the paperwork for its listing public, so details of their accounts are sketchy. But Anthropic is getting the better of most comparisons.

CHART: THE ECONOMIST

Because the two firms are both young and fast-growing, they measure their notional annual turnover by multiplying the figure for the most recent month by 12—a number called annual recurring revenue (ARR). Futurum, a tech-research firm, notes that Anthropic accounts differently for sales through cloud-computing partners, which may flatter its figures. Nonetheless, the trend is unmistakable. Anthropic's ARR has raced past OpenAI's (see chart). Anthropic's valuation, too, has surpassed OpenAI's, reaching $965bn in May, compared with OpenAI's $852bn in March. Anthropic has told investors it will generate an operating profit in the second quarter, which would make it the first AI lab to do so, according to Futurum.

There are other signs of strength. For all the talk of an AI bubble, Anthropic's most recent valuation, at just over 20 times ARR, should be judged in the context of recurring revenues that grew fivefold in the past five months. "It is not necessarily cheap but at these growth rates it is not insane," says Harrison Rolfes of PitchBook, a data-gatherer. OpenAI is more richly valued. The recent IPO of SpaceX valued it at more than 90 times revenues.

But Anthropic's lead is not insurmountable. OpenAI is reportedly considering launching a price war, dramatically cutting what it charges per token (the chunks of data processed by AI) in response to concerns from businesses that they are spending too much on AI. Although the price of tokens is already falling, Anthropic's latest models use more tokens per inquiry. Anthropic's pricing power means it did not need to offer a commensurate discount, raising the price of each request. This may have helped fuel its revenue growth.

More generally, AI models are at risk of commoditisation because rival firms tend to catch up with each other quickly. That requires model-makers to expand into other domains, especially through the use of agents. Anthropic has been the main source of concerns in the software-as-a-service industry that AI labs will supplant SaaS companies providing tools for things like procurement and personnel management. That would not only give it a much stronger relationship with its customers, but also gain it access to more data that could be used to improve its models' performance at such tasks. OpenAI has its eye on the same prize, as does Google, whose enterprise-software products are already deeply embedded in many firms. But for now, at least, Anthropic's business focus gives it an edge.

All the while, Mr Amodei continues to preach about the risks of AI—literally. Every two weeks or so, he delivers hour-long talks to Anthropic's 3,000 employees. The firm's day-to-day operations are led by Ms Amodei; she too shares her brother's convictions about the importance of AI safety. Further down the ranks, the determination to be the first firm to create artificial general intelligence (AGI), or superhuman AI, runs deep. AI researchers at other firms note how "AGI-pilled" Anthropic's employees are, despite the conviction among many of them that its advent will put them out of work.

It may have been all the hand-wringing about the potential downsides of AI that initially irked the Trump administration. Officials have blamed the "doomer industrial complex" at Anthropic for what they see as excessive regulation of the technology during the administration of Joe Biden (some members of which have since joined the firm). For much of Mr Trump's term, the government has taken a dim view of attempts to slow down AI, focusing instead on staying ahead of China. Last year, while some of Anthropic's rivals, such as OpenAI, openly embraced Mr Trump's "accelerationist" approach, Mr Amodei was among the loudest voices calling for more regulation at the state and federal level.

That did not go down well. David Sacks, an adviser to the government on AI, accused Anthropic of running a "sophisticated regulatory-capture strategy based on fearmongering". Other officials noted the irony of warning about AI while building ever more capable models.

Turbulent priest

At any rate, it was Anthropic's tendency to highlight the dangers of AI that seems to have led to its first public falling-out with the Trump administration, in late February. The firm and the Department of War got into a public row about the use of Claude to power mass domestic surveillance or autonomous weapons. The firm said that existing laws were not adequate to regulate the use of AI for mass surveillance and that AI was not reliable enough to be used in fully autonomous killing machines. The Pentagon suggested that Mr Amodei was trying to put himself in charge of matters of national security. In the end Mr Trump sought to compel all government agencies to sever ties with Anthropic, whose leaders he denounced as "leftwing nut jobs".

The incident appears to have stoked a fear within the administration that Anthropic is becoming so successful that it may be a threat in itself. Mr Sacks recently warned that if Anthropic's "exponential" growth continues for long, it could gain "unprecedented control over the most important technology of our time". He drew an unflattering parallel with John D. Rockefeller's domination of the oil industry in the 19th century.

In particular, members of the administration have woken up to the massive risks high-powered AI poses to national security. That seems to have been the origin of the latest spat between Anthropic and the authorities. The firm had been co-operating with the government over the release of its first Mythos model, which is superbly good at hacking, among other alarming capabilities, and had found potential vulnerabilities in many government systems. Anthropic allowed a limited number of customers, vetted by the government, to use it. Later it created a version of Mythos called Fable, with special guardrails to prevent abuse, which it distributed more widely.

Some reports suggest that Anthropic was slow to tell the government when it gave new customers access to Mythos. A source close to Anthropic disputes this, saying it worked closely with the administration, allowing it to raise national-security concerns about any new partners and swiftly cutting off a global telecommunications firm when the government asked it to. (In one of many dark insinuations about Anthropic, the Trump administration has whispered that the customer in question had links to China. That may be true, but it is not a sign of hidden sympathies on Anthropic's part. Few CEOs are as loudly critical of Chinese authoritarianism as Mr Amodei, who has strongly advocated strengthening America's curbs on AI chip exports to China.)

Matters came to a head after Amazon, the e-commerce giant that is a big investor in Anthropic, told the authorities that it had found security flaws in Fable, enabling it to "jailbreak", or override the guardrails. Posting on X, Mr Sacks accused Mr Amodei of refusing to "fix the jailbreak".

Sacks accused Mr Amodei of refusing to "fix the jailbreak".

Confusion at Canossa

A source close to Anthropic says there was no refusal. The firm insists it gave the models to the government to probe for flaws weeks before Fable was released. It also disputes how serious the jailbreak was. There is no such thing, it has argued, as completely foolproof security controls for an AI model. Many observers assume that Anthropic was penalised chiefly out of personal animosity. After all, in February the administration accused Mr Amodei of having a "God complex" and Mr Amodei said Anthropic was being persecuted for failing to offer Mr Trump "dictator-style praise".

But underpinning such squabbling is a genuine regulatory conundrum. Having adopted a laissez-faire approach to AI, the Trump administration does not have many straightforward or calibrated ways of imposing its will on AI labs. Adam Thierer of the R Street Institute, a think-tank, says Mythos set off a tussle within the government over whether models had come to pose a sufficient threat to national security to require vetting before release. Hawkish types called for an official say over the release of new models. Some even floated the creation of a dedicated regulatory agency to that end. Instead, earlier this month Mr Trump issued an executive order creating a voluntary framework through which AI labs could work with the government before releasing new models. It is not yet up and running.

ILLUSTRATION: TIMO LENZEN

Without regulations in place, the administration is left with blunt tools. During Anthropic's row with the Pentagon, officials ended up declaring the firm a "supply-chain risk", thereby impeding government agencies from doing business with it—a potentially self-defeating stance to take towards America's leading AI lab. In the ongoing row over Mythos and Fable, the administration is said to have required Anthropic to seek a special export licence for every sale to a foreign entity. Since Anthropic cannot be sure whether foreigners working at American firms, for instance, might be using the models, the company ended up having to shut them down altogether—again, an outcome likely to do America more harm than good. Without clear rules, says Dean Ball, a former

America more harm than good. Without clear rules, says Dean Ball, a former Trump administration official at the Foundation for American Innovation, a think-tank, the risk is that America winds up with government “dominance over AI” rather than American “dominance in AI”.

The tensions are natural. It makes sense for a firm like Anthropic to want to distribute a powerful model as widely as possible. It is good for business because it caters to a burning need, in that criminals are already using AI to make their cyber-attacks more effective. It is equally understandable for the national-security apparatus to want to take charge of such a powerful technology, both for offensive and defensive purposes. “The national-security apparatus is in control mode,” Mr Ball says.

The question is how control of the technology will be exercised. Although Mr Amodei sometimes talks in alarmist terms about the risks of AI, his favoured approach is a narrowly technocratic form of regulation. He calls for rules modelled on regulatory agencies such as the Federal Aviation Administration, whereby frontier AI models would be put through standardised technical tests and audits.

Yet even Mr Amodei acknowledges that, as models become ever more capable, AI will become the “dominant source of military and economic power”. That gives the Trump administration and all its successors a clear and compelling interest in having their own levers to pull, including over Anthropic. As Mr Ball puts it, “No company gets to shake the foundations of state sovereignty while staying blithely above the raw reality of politics.”

Both sides hope the row will blow over. It has benefited neither of them. In a capricious attempt to assert its authority, the government has upset people around the world who use Anthropic’s models, as well as showing how short-sighted it was to dismiss the need for rules and regulations on AI in the first place. For its part Anthropic, acting like it knows best rather than bending over backwards to accommodate an antagonistic government, has ended up hamstrung, to the detriment of customers.

The fallout will hang over Anthropic’s IPO. The suspension of a commercial model deployed to hundreds of millions of people is bound to affect its bottom line. What is more, investors and customers alike will worry about the risk of further such disruption owing to Anthropic’s testy relationship with the authorities.

Ultimately, AI firms, for all the good they may do, are developing dangerous capabilities at a time when the laws to govern them are inadequate. Mr Amodei knows that regulations evolve far more slowly than models, so he finds himself acting as an unelected arbiter of safety. The elected government sees AI as too potent a technology to pass up. The tussle will persist until Congress writes new rules. ■

Explore more

Artificial intelligence — Follow

Briefing — Follow

United States — Follow

Anthropic — Follow

→ See the latest from topics you follow

This article appeared in the Briefing section of the print edition under the headline “A test of faith”

From the June 20th 2026 edition

Discover stories from this section and more in the list of contents

Explore the edition

Save Share Reuse this content

More from Briefing →

Donald Trump's blocking of Anthropic is capricious and chaotic

America's closest allies are shellshocked

Forget the World Cup. Culture is becoming more fragmented

Farewell to the monoculture

India's population will soon be falling—probably quite fast

Neither widespread poverty, nor high rates of marriage nor relatively young mothers are sustaining fertility

Elon Musk is going all-in on an unproven technology

What lies behind the world's biggest public offering

Indonesia’s president is jeopardising the economy and democracy

Prabowo Subianto is too spendthrift and too authoritarian

China is pushing Donald Trump for Taiwan concessions

Some in Taipei and Washington fear he may give ground